

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: data = pd.read_csv(r'StudentsPerformance.csv')
```

```
In [4]: data
```

```
Out[4]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group B	bachelor's degree	standard	none	72	72	74
1	female	group C	some college	standard	completed	69	90	88
2	female	group B	master's degree	standard	none	90	95	93
3	male	group A	associate's degree	free/reduced	none	47	57	44
4	male	group C	some college	standard	none	76	78	75
...
995	female	group E	master's degree	standard	completed	88	99	95
996	male	group C	high school	free/reduced	none	62	55	55
997	female	group C	high school	free/reduced	completed	59	71	65
998	female	group D	some college	standard	completed	68	78	77
999	female	group D	some college	free/reduced	none	77	86	86

1000 rows × 8 columns

```
In [5]: # top five rows
data.head()
```

Out[5]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group B	bachelor's degree	standard	none	72	72	74
1	female	group C	some college	standard	completed	69	90	88
2	female	group B	master's degree	standard	none	90	95	93
3	male	group A	associate's degree	free/reduced	none	47	57	44
4	male	group C	some college	standard	none	76	78	75

In [6]: `# to check data information`
`data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                1000 non-null   object
1   race/ethnicity                        1000 non-null   object
2   parental level of education           1000 non-null   object
3   lunch                                 1000 non-null   object
4   test preparation course               1000 non-null   object
5   math score                           1000 non-null   int64
6   reading score                        1000 non-null   int64
7   writing score                         1000 non-null   int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB
```

In [7]: `data['gender'].dtypes`

Out[7]: `dtype('O')`

In [8]: `data['gender'].dtypes=='O'`

Out[8]: `True`

In [9]: `data.columns`

Out[9]: `Index(['gender', 'race/ethnicity', 'parental level of education', 'lunch', 'test preparation course', 'math score', 'reading score', 'writing score'], dtype='object')`

In [10]: `# iterate through columns`
`[i for i in data.columns]`

```
Out[10]: ['gender',
          'race/ethnicity',
          'parental level of education',
          'lunch',
          'test preparation course',
          'math score',
          'reading score',
          'writing score']
```

```
In [11]: # pull out categorical columns
cat_col = [i for i in data.columns if data[i].dtype == 'O']
cat_col
```

```
Out[11]: ['gender',
          'race/ethnicity',
          'parental level of education',
          'lunch',
          'test preparation course']
```

```
In [12]: # pull out numerical columns
num_col = [i for i in data.columns if data[i].dtype != 'O']
num_col
```

```
Out[12]: ['math score', 'reading score', 'writing score']
```

```
In [13]: data[num_col]
```

```
Out[13]:
```

	math score	reading score	writing score
0	72	72	74
1	69	90	88
2	90	95	93
3	47	57	44
4	76	78	75
...
995	88	99	95
996	62	55	55
997	59	71	65
998	68	78	77
999	77	86	86

1000 rows × 3 columns

```
In [14]: data[cat_col]
```

Out[14]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course
0	female	group B	bachelor's degree	standard	none
1	female	group C	some college	standard	completed
2	female	group B	master's degree	standard	none
3	male	group A	associate's degree	free/reduced	none
4	male	group C	some college	standard	none
...
995	female	group E	master's degree	standard	completed
996	male	group C	high school	free/reduced	none
997	female	group C	high school	free/reduced	completed
998	female	group D	some college	standard	completed
999	female	group D	some college	free/reduced	none

1000 rows × 5 columns

In [15]: *# to check memory usage of this dataset*
`data.memory_usage()`

Out[15]:

Index	128
gender	8000
race/ethnicity	8000
parental level of education	8000
lunch	8000
test preparation course	8000
math score	8000
reading score	8000
writing score	8000
dtype: int64	

Missing value

In [16]: `data.isnull()`

Out[16]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
995	False	False	False	False	False	False	False	False
996	False	False	False	False	False	False	False	False
997	False	False	False	False	False	False	False	False
998	False	False	False	False	False	False	False	False
999	False	False	False	False	False	False	False	False

1000 rows × 8 columns

```
In [17]: # sum of individual column
data.isnull().sum()
```

```
Out[17]: gender                0
race/ethnicity              0
parental level of education  0
lunch                      0
test preparation course      0
math score                  0
reading score                0
writing score                0
dtype: int64
```

```
In [18]: # sum of all
data.isnull().sum().sum()
```

```
Out[18]: 0
```

```
In [19]: # to check duplicate data
data.duplicated()
```

```
Out[19]: 0      False
1      False
2      False
3      False
4      False
...
995    False
996    False
997    False
998    False
999    False
Length: 1000, dtype: bool
```

```
In [20]: # to check sum of duplicate values
data.duplicated().sum()
```

Out[20]: 0

In [21]: *# to check unique value column wise*
data.nunique()

Out[21]:

gender	2
race/ethnicity	5
parental level of education	6
lunch	2
test preparation course	2
math score	81
reading score	72
writing score	77
dtype:	int64

In [22]: *# sum of unique values*
data.nunique().sum()

Out[22]: 247

In [23]: *# to check unique value of a column*
data['gender'].unique()

Out[23]: array(['female', 'male'], dtype=object)

In [24]: *# to check unique value of a column*
data['race/ethnicity'].unique()

Out[24]: array(['group B', 'group C', 'group A', 'group D', 'group E'],
dtype=object)

In [25]: *# statistic of the dataset*
data.describe().T

Out[25]:

	count	mean	std	min	25%	50%	75%	max
math score	1000.0	66.089	15.163080	0.0	57.00	66.0	77.0	100.0
reading score	1000.0	69.169	14.600192	17.0	59.00	70.0	79.0	100.0
writing score	1000.0	68.054	15.195657	10.0	57.75	69.0	79.0	100.0

In [26]: *# to check correlation*
data.corr()

Out[26]:

	math score	reading score	writing score
math score	1.000000	0.817580	0.802642
reading score	0.817580	1.000000	0.954598
writing score	0.802642	0.954598	1.000000


In [27]: *# to check skewness of the data*
data.skew()

Out[27]:

math score	-0.278935
reading score	-0.259105
writing score	-0.289444
dtype:	float64


```
In [28]: #to check distribution plot for a lleft skewed column  
sns.distplot(data['math score'])
```

```
Out[28]: <AxesSubplot:xlabel='math score', ylabel='Density'>
```




```
In [29]: sns.displot(data['math score'],kde=True)
```

```
Out[29]: <seaborn.axisgrid.FacetGrid at 0x20dc1af6ee0>
```



```
In [30]: sns.kdeplot(data['math score'])
```

```
Out[30]: <AxesSubplot:xlabel='math score', ylabel='Density'>
```



```
In [31]: data
```


Out[31]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group B	bachelor's degree	standard	none	72	72	74
1	female	group C	some college	standard	completed	69	90	88
2	female	group B	master's degree	standard	none	90	95	93
3	male	group A	associate's degree	free/reduced	none	47	57	44
4	male	group C	some college	standard	none	76	78	75
...
995	female	group E	master's degree	standard	completed	88	99	95
996	male	group C	high school	free/reduced	none	62	55	55
997	female	group C	high school	free/reduced	completed	59	71	65
998	female	group D	some college	standard	completed	68	78	77
999	female	group D	some college	free/reduced	none	77	86	86

1000 rows × 8 columns

In [32]: `data.columns`

Out[32]: Index(['gender', 'race/ethnicity', 'parental level of education', 'lunch', 'test preparation course', 'math score', 'reading score', 'writing score'], dtype='object')

In [33]: `data['math score'] + data['reading score'] + data['writing score']`

Out[33]:

0	218
1	247
2	278
3	148
4	229
...	
995	282
996	172
997	195
998	223
999	249

Length: 1000, dtype: int64

In [38]: `# adding a new column 'average'`
`data['average'] = (data['math score'] + data['reading score'] + data['writing score'])`

In [39]: `data.head()`

Out[39]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
0	female	group B	bachelor's degree	standard	none	72	72	74	72.666667
1	female	group C	some college	standard	completed	69	90	88	82.333333
2	female	group B	master's degree	standard	none	90	95	93	92.666667
3	male	group A	associate's degree	free/reduced	none	47	57	44	49.333333
4	male	group C	some college	standard	none	76	78	75	76.333333



In [41]: `data.groupby('gender').mean()`

Out[41]:

	math score	reading score	writing score	average
gender				
female	63.633205	72.608108	72.467181	69.569498
male	68.728216	65.473029	63.311203	65.837483

In [42]: `data.groupby('gender').count()`

Out[42]:

	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
gender								
female	518	518	518	518	518	518	518	518
male	482	482	482	482	482	482	482	482

In [48]: `#who get less than 30 in math`
`data[data['math score']<30]`

Out[48]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
17	female	group B	some high school	free/reduced	none	18	32	28	26.0000
59	female	group C	some high school	free/reduced	none	0	17	10	9.0000
91	male	group C	high school	free/reduced	none	27	34	36	32.3333
145	female	group C	some college	free/reduced	none	22	39	33	31.3333
327	male	group A	some college	free/reduced	none	28	23	19	23.3333
338	female	group B	some high school	free/reduced	none	24	38	27	29.6666
363	female	group D	some high school	free/reduced	none	27	34	32	31.0000
466	female	group D	associate's degree	free/reduced	none	26	31	38	31.6666
528	female	group D	bachelor's degree	free/reduced	none	29	41	47	39.0000
601	female	group C	high school	standard	none	29	29	30	29.3333
683	female	group C	some high school	free/reduced	completed	29	40	44	37.6666
787	female	group B	some college	standard	none	19	38	32	29.6666
842	female	group B	high school	free/reduced	completed	23	44	36	34.3333
980	female	group B	high school	free/reduced	none	8	24	23	18.3333

In [49]: `data[data['math score']<30].count()`

Out[49]:

gender	14
race/ethnicity	14
parental level of education	14
lunch	14
test preparation course	14
math score	14
reading score	14
writing score	14
average	14
dtype:	int64

In [51]:

```
data_num = data[num_col]
data_num
```

Out[51]:

	math score	reading score	writing score
0	72	72	74
1	69	90	88
2	90	95	93
3	47	57	44
4	76	78	75
...
995	88	99	95
996	62	55	55
997	59	71	65
998	68	78	77
999	77	86	86

1000 rows × 3 columns

In [52]: `from scipy.stats import normaltest`

In [55]: `normaltest(data_num['math score'])`

Out[55]: NormaltestResult(statistic=15.408960513931822, pvalue=0.00045080293869937836)


In [58]: `normaltest(data_num['math score'])[1]*100`

Out[58]: 0.04508029386993784

In []: *# if p value > 0.05 then my data will be normally distributed*
if p value < 0.05 then data will be non normally distributed


In [61]: `sns.distplot(data_num['math score'])`

Out[61]: <AxesSubplot:xlabel='math score', ylabel='Density'>




```
In [71]: sns.distplot(data['writing score'])
```

```
Out[71]: <AxesSubplot:xlabel='writing score', ylabel='Density'>
```




```
In [72]: sns.distplot(data['reading score'])
```

```
Out[72]: <AxesSubplot:xlabel='reading score', ylabel='Density'>
```




```
In [63]: # outliers  
sns.boxplot(data=data['math score'])
```

Out[63]: <AxesSubplot:>




```
In [64]: sns.boxplot(data['math score'])
```

Out[64]: <AxesSubplot:xlabel='math score'>




```
In [65]: sns.boxplot(data=data['reading score'])
```

```
Out[65]: <AxesSubplot:>
```



```
In [66]: sns.boxplot(data=data['reading score'])
```

```
Out[66]: <AxesSubplot:>
```



```
In [92]: q1 = data['math score'].quantile(0.25)
q1
```

```
Out[92]: 57.0
```

```
In [93]: q3 = data['math score'].quantile(0.75)
q3
```

```
Out[93]: 77.0
```

```
In [94]: #Interquartile range
IQR = q3 - q1
IQR
```

```
Out[94]: 20.0
```

```
In [95]: upper_limit = q3 + (1.5*IQR)
upper_limit
```

```
Out[95]: 107.0
```

```
In [96]: lower_limit = q1 - (1.5*IQR)
lower_limit
```

```
Out[96]: 27.0
```

```
In [97]: data['math score'].min()
```

```
Out[97]: 0
```

```
In [98]: data['math score'].max()
```

```
Out[98]: 100
```


In [99]: `data['math score'].unique()`

Out[99]: `array([72, 69, 90, 47, 76, 71, 88, 40, 64, 38, 58, 65, 78,
50, 18, 46, 54, 66, 44, 74, 73, 67, 70, 62, 63, 56,
97, 81, 75, 57, 55, 53, 59, 82, 77, 33, 52, 0, 79,
39, 45, 60, 61, 41, 49, 30, 80, 42, 27, 43, 68, 85,
98, 87, 51, 99, 84, 91, 83, 89, 22, 100, 96, 94, 48,
35, 34, 86, 92, 37, 28, 24, 26, 95, 36, 29, 32, 93,
19, 23, 8], dtype=int64)`

In [100... `# to check outlier in math score column which has less than lower limit 27
data[data['math score'] < lower_limit]`

Out[100]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
17	female	group B	some high school	free/reduced	none	18	32	28	26.0000
59	female	group C	some high school	free/reduced	none	0	17	10	9.0000
145	female	group C	some college	free/reduced	none	22	39	33	31.3333
338	female	group B	some high school	free/reduced	none	24	38	27	29.6666
466	female	group D	associate's degree	free/reduced	none	26	31	38	31.6666
787	female	group B	some college	standard	none	19	38	32	29.6666
842	female	group B	high school	free/reduced	completed	23	44	36	34.3333
980	female	group B	high school	free/reduced	none	8	24	23	18.3333

In [102... `data[data['math score'] > upper_limit]`

Out[102]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
--	--------	----------------	-----------------------------------	-------	-------------------------------	---------------	------------------	------------------	---------

In [103... `def outlier_threshold(df, variable):
 q1 = df[variable].quantile(0.25)
 q3 = df[variable].quantile(0.75)
 iqr = q3 - q1
 upper_limit = q3 + (1.5*iqr)
 lower_limit = q1 - (1.5*iqr)
 return upper_limit, lower_limit`

In [104... `def get_iqr(df, column_name, variable, q1_range, q3_range):
 q1 = df[column_name].quantile(q1_range)
 q3 = df[column_name].quantile(q3_range)
 IQR = q3 - q1
 upper_fence = q3 + (1.5*IQR)
 lower_fence = q1 - (1.5*IQR)
 return IQR, upper_fence, lower_fence`

In [112... data_num.columns

Out[112]: Index(['math score', 'reading score', 'writing score'], dtype='object')

```
In [111... # to check upper_limit, lower_limit for each columns
for variable in data_num.columns:
    upper_limit, lower_limit = outlier_threshold(data_num, variable)
    print(upper_limit, lower_limit)
```

107.0 27.0
109.0 29.0
110.875 25.875

```
In [ ]: def replace_with_threshold(data, numeric_col):
    for variable in numeric_col:
        upper_limit, lower_limit = outlier_threshold(data_num, variable)
        data.loc[data[variable]<lower_limit, variable] = lower_limit
        data.loc[data[variable]>upper_limit, variable] = upper_limit
```

```
In [128... #to check how many data points are lower than lower fence or lower limit
data.loc[data['math score']<lower_limit, 'math score']
```

Out[128]: 17 18
59 0
145 22
338 24
787 19
842 23
980 8
Name: math score, dtype: int64

In [133... data[data['math score']<lower_limit]

Out[133]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	avera
17	female	group B	some high school	free/reduced	none	18	32	28	26.0000
59	female	group C	some high school	free/reduced	none	0	17	10	9.0000
145	female	group C	some college	free/reduced	none	22	39	33	31.3333
338	female	group B	some high school	free/reduced	none	24	38	27	29.6666
787	female	group B	some college	standard	none	19	38	32	29.6666
842	female	group B	high school	free/reduced	completed	23	44	36	34.3333
980	female	group B	high school	free/reduced	none	8	24	23	18.3333

```
In [129... #to check how many data points are higher than upper fence or upper limit
data.loc[data['math score']>upper_limit, 'math score']
```

Out[129]: Series([], Name: math score, dtype: int64)

In [130... data['math score'].max()

Out[130]: 100

In [134... data[data['math score']>upper_limit]

Out[134]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
--	--------	----------------	-----------------------------------	-------	-------------------------------	---------------	------------------	------------------	---------

In [118... data.head()

Out[118]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
0	female	group B	bachelor's degree	standard	none	72	72	74	72.666667
1	female	group C	some college	standard	completed	69	90	88	82.333333
2	female	group B	master's degree	standard	none	90	95	93	92.666667
3	male	group A	associate's degree	free/reduced	none	47	57	44	49.333333
4	male	group C	some college	standard	none	76	78	75	76.333333



In [135... data

Out[135]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
0	female	group B	bachelor's degree	standard	none	72	72	74	72.66666666666667
1	female	group C	some college	standard	completed	69	90	88	82.33333333333333
2	female	group B	master's degree	standard	none	90	95	93	92.66666666666667
3	male	group A	associate's degree	free/reduced	none	47	57	44	49.33333333333333
4	male	group C	some college	standard	none	76	78	75	76.33333333333333
...
995	female	group E	master's degree	standard	completed	88	99	95	94.00000000000001
996	male	group C	high school	free/reduced	none	62	55	55	57.33333333333333
997	female	group C	high school	free/reduced	completed	59	71	65	65.00000000000001
998	female	group D	some college	standard	completed	68	78	77	74.33333333333333
999	female	group D	some college	free/reduced	none	77	86	86	83.00000000000001

1000 rows × 9 columns




Graph Analysis

In [136]:

```
sns.countplot(data['gender'])
```


Out[136]:

```
<AxesSubplot:xlabel='gender', ylabel='count'>
```



```
In [139... sns.countplot(data['race/ethnicity'])
```

```
Out[139]: <AxesSubplot:xlabel='race/ethnicity', ylabel='count'>
```



```
In [138... data.head()
```

Out[138]:

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score	average
0	female	group B	bachelor's degree	standard	none	72	72	74	72.666667
1	female	group C	some college	standard	completed	69	90	88	82.333333
2	female	group B	master's degree	standard	none	90	95	93	92.666667
3	male	group A	associate's degree	free/reduced	none	47	57	44	49.333333
4	male	group C	some college	standard	none	76	78	75	76.333333

In [141... `df = data.groupby('gender').mean()`
df

Out[141]:

	math score	reading score	writing score	average
gender				

female	63.633205	72.608108	72.467181	69.569498
male	68.728216	65.473029	63.311203	65.837483

In [142... `df['average']`

Out[142]:

```
gender
female    69.569498
male      65.837483
Name: average, dtype: float64
```

In [143... `df['average'][0]`

Out[143]: 69.56949806949807

In [145... `df['average'][1]`

Out[145]: 65.8374827109267

In [146... `df['math score'][0]`

Out[146]: 63.633204633204635

In [147... `df['math score'][1]`

Out[147]: 68.72821576763485

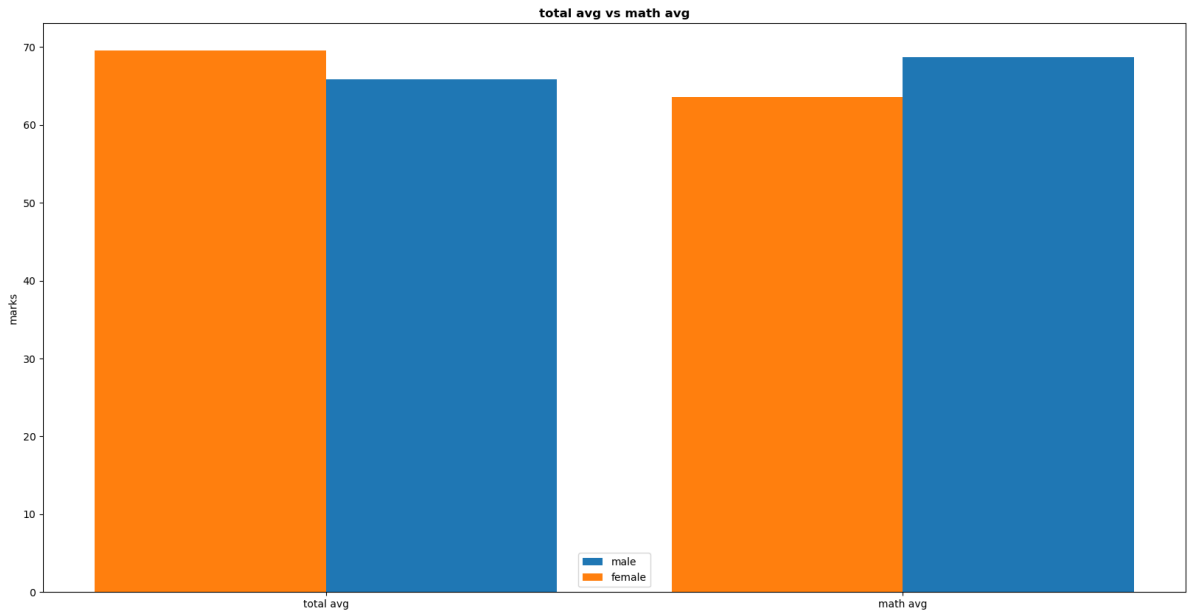
In [149... `plt.figure(figsize = (20,10))`
`x = ['total avg', 'math avg']`
`female_score = df['average'][0], df['math score'][0]`
`male_score = df['average'][1], df['math score'][1]`

`x_axis = np.arange(len(x))`
`plt.bar(x_axis + 0.2, male_score, 0.4, label = 'male')`

```
plt.bar(x_axis - 0.2, female_score, 0.4, label = 'female')


plt.xticks(x_axis,x)
plt.ylabel('marks')
plt.title('total avg vs math avg', fontweight = 'bold')
plt.legend()
plt.show
```

Out[149]: <function matplotlib.pyplot.show(close=None, block=None)>



In [150... sns.heatmap(df)

Out[150]: <AxesSubplot:ylabel='gender'>




In [163... data_num.corr()

Out[163]:

	math score	reading score	writing score
math score	1.000000	0.817580	0.802642
reading score	0.817580	1.000000	0.954598
writing score	0.802642	0.954598	1.000000

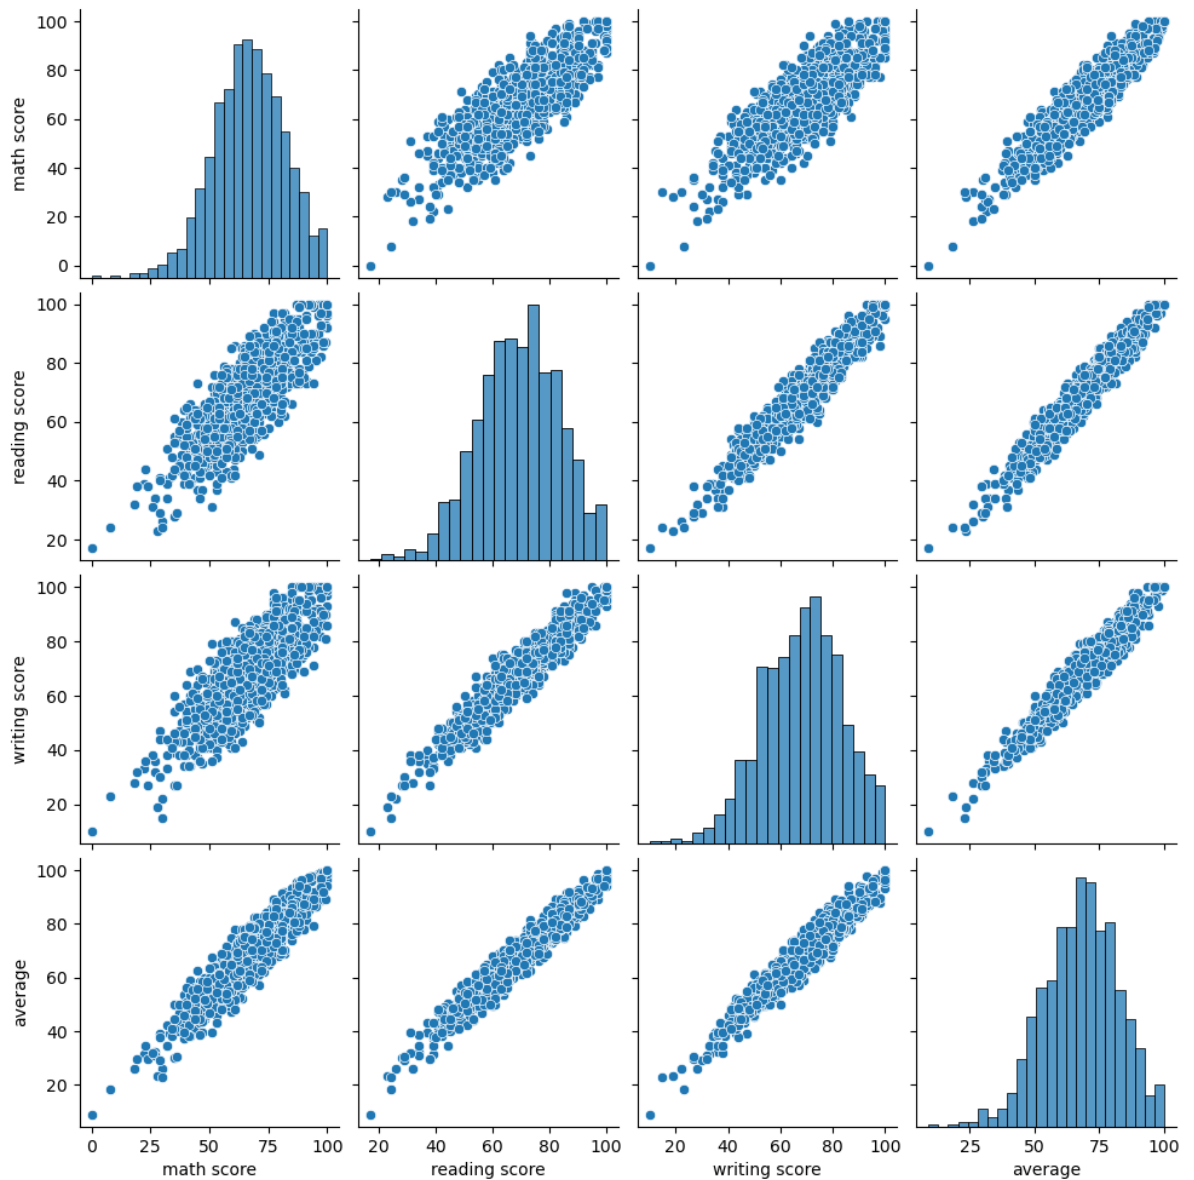
In [158... `sns.heatmap(data_num.corr(), annot = True)`Out[158]: `<AxesSubplot:>`

```
In [162... sns.heatmap(data_num.corr(), annot = True, cmap = 'icefire', linewidths = 0.3)
fig = plt.gcf()
fig.set_size_inches(12,8)
plt.title('Correlation between variables', color = 'black', size = 15)
plt.show()
```


```
In [157... sns.pairplot(data)
```

```
Out[157]: <seaborn.axisgrid.PairGrid at 0x20dce2a0790>
```



```
In [167... sns.violinplot(data = data_num)
```

```
Out[167]: <AxesSubplot:>
```



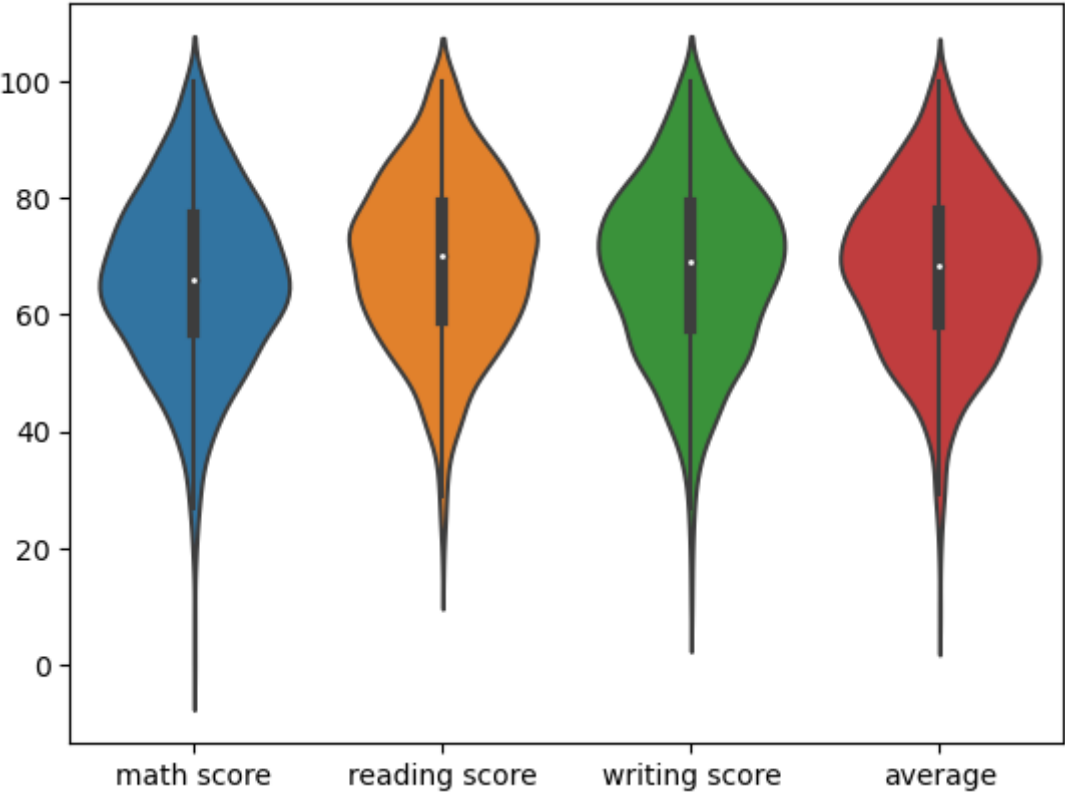
```
In [164... df.head()
```

Out[164]:

	math score	reading score	writing score	average
gender				
female	63.633205	72.608108	72.467181	69.569498
male	68.728216	65.473029	63.311203	65.837483

```
In [168... sns.violinplot(data = data)
```

Out[168]: <AxesSubplot:>



```
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
```